

Biomedical Statistics Coursework2

Junrui Wang

May 3, 2022

Contents

1	Intro	2
2	Data Display	2
3	Reported COVID-19 deaths	2
3.1	Extract the reported COVID-19 deaths	2
3.2	Aggregate the reported COVID-19 deaths	2
3.3	Plot the monthly, reported COVID-19 deaths in the city	3
3.4	Calculate the cumulating, monthly reported COVID-19 deaths in the city that you are assigned, again by gender and age group.	4
3.5	Plot for the cumulating, monthly reported COVID-19 deaths and the cumulating monthly excess deaths for each gender and age group in the city	5
3.6	Calculate the cumulating, monthly reported COVID-19 deaths in the city that you are assigned, per 1,000 population in each gender and age group.	5
3.7	Make a suitable plot that shows the total reported COVID-19 deaths by the end of 2021 across all cities, per 1,000 population in each city. Show cities on the x-axis and the total per 1,000 population on the y-axis.	6
3.8	Comment if this plot is a fair comparison of total reported COVID-19 deaths across cities.	7
4	Assessing the Coursework 1 model (6 points)	7
4.1	Assess convergence, mixing, and make the trace plot corresponding to the model parameter with smallest effective sample size.	10
4.2	Make a posterior predictive check and report the proportion of data points inside the 95% posterior predictive credibility intervals associated with each data point, separately for men and women. . . .	11
4.3	Comment if the model from Coursework 1 provides a sufficiently good fit to the data for the city that you are assigned.	12
5	Poisson model with random time effects	13
5.1	Describe Bayesian Poisson model	13
5.2	Implement the hierarchical Bayesian Poisson model in Stan.	13
5.3	Check convergence and mixing	15
5.4	Make a posterior predictive check and report the proportion of data points inside the 95% posterior predictive credibility intervals associated with each data point, separately for men and women. . . .	16
5.5	Plot the posterior predictive checks	17
5.6	Comment	18
6	Your best model	18
6.1	Implement your hierarchical Bayesian Poisson model in Stan.	18
6.2	Check convergence and mixing in an appropriate manner	20
6.3	Quantify model fit using a posterior predictive check.	21

7	Under-reporting of COVID-19 deaths	22
7.1	Calculate the monthly excess deaths for each gender and age group	24
7.2	Plot	24
7.3	Calculate the difference in total excess deaths and the reported COVID-19 deaths by the end of 2021 in the city that you are assigned. Report posterior median estimates along with 95% posterior credible intervals for both men and women in a table.	25
8	Evaluating prediction accuracy	26
8.1	26
8.2	28
8.3	Rio de Janeiro	29
8.4	Comment on the within-sample and out-of-sample prediction accuracies of your best model.	31

1 Intro

"I, Junrui Wang, certify that this assessed coursework is my own and unaided work, unless otherwise acknowledged through citations and references. I have not discussed my coursework with anyone else except when seeking clarification with the module lecturer via email or on MS Teams. I understand that all rules regarding academic integrity and plagiarism apply, and that violations of this will be treated as an examination offence. In particular, I have not shared any of my work with anyone else prior to submission."

2 Data Display

3 Reported COVID-19 deaths

The city allocated to me is Belo Horizonte.

3.1 Extract the reported COVID-19 deaths

```

1 data.dir <- "/ic math/MSc Stats Data Science/Bio Stats/coursework/Coursework2"
2 out.dir <- "/ic math/MSc Stats Data Science/Bio Stats/coursework/Coursework2/plots"
3 dd <- read.csv(file = file.path(data.dir, "civil_registry_covid_cities_detailed.csv")) %>%
4   as.data.table()
5 mycolumns <- c('date', 'city', 'place', 'gender', 'age_group',
6   'deaths_covid19', 'deaths_stroke_covid19', 'deaths_heart_attack_covid19')
7 dfcd <- dd[, mycolumns, with = FALSE]
8 dfcd <- melt(dfcd, id.vars = c('date', 'city', 'place', 'gender', 'age_group'))
9 set(dfcd, dfcd[, which(is.na(value))], 'value', 0L)
10 dfcd <- dfcd[, list(deaths = sum(value)), by = c('date', 'city', 'gender', 'age_group')]

```

The first 5 data in the extracted dataset is

	date	city	gender	age_group	deaths
1	2019-01-01	Rio Branco	F	80-89	0
2	2019-01-01	Rio Branco	M	30-39	0
3	2019-01-01	Rio Branco	M	50-59	0
4	2019-01-01	Rio Branco	M	80-89	0
5	2019-01-01	Rio Branco	M	9-	0

3.2 Aggregate the reported COVID-19 deaths

```

1 dfcd[, month := as.integer(substr(date,6,7 ))]
2 dfcd[, year := as.integer(substr(date,1,4))]
3 dfcd[, month_id := (year - 2019) * 12 + month]
4
5 # specify age groups as requested
6 set(dfcd, dfcd[, which(age_group == '9-')], 'age_group', '0-29')
7 set(dfcd, dfcd[, which(age_group == '10-19')], 'age_group', '0-29')
8 set(dfcd, dfcd[, which(age_group == '20-29')], 'age_group', '0-29')
9 set(dfcd, dfcd[, which(age_group == '90-99')], 'age_group', '>=90')
10 set(dfcd, dfcd[, which(age_group == '100+')], 'age_group', '>=90')
11 set(dfcd, dfcd[, which(is.na(age_group))], 'age_group', 'unknown')
12 dfcd <- subset(dfcd, ! age_group %in% c('>=90','unknown'))
13
14 # Group
15 cvid_agg <- dfcd[,.(Deaths = sum(deaths)), by = .(gender,age_group, month_id)]
16 setkey(cvid_agg, gender, age_group, month_id)

```

Group by gender, the age groups 0-29, 30-39, 40-49, 50-59, 60-69, 70-79, 80-89 (ignoring deaths in older ages), and by continuous month index we get the following table

	gender	age_group	month_id	Deaths
1	F	0-29	1.00	0
2	F	0-29	2.00	0
3	F	0-29	3.00	0
4	F	0-29	4.00	0
5	F	0-29	5.00	0

Table 1: The first 5 elements from *cvid_agg*

3.3 Plot the monthly, reported COVID-19 deaths in the city

```

1 mycity <- dfcd[which(city == "Belo Horizonte"),]
2
3 mycity_agg <- mycity[,.(Deaths = sum(deaths)), by = .(gender,age_group, month_id)]
4 p <- ggplot(mycity_agg , aes(x = month_id, y = Deaths, fill = gender)) +
5   geom_col() + facet_grid(~age_group) +
6   theme_bw()+
7   labs( x='Time/Month', y='Monthly Deaths') +
8   theme(axis.text.x =element_text(size = 6))
9
10 ggsave(file=file.path(out.dir, 'monthly_death.png'), p, w=5, h=5)
11 include_graphics(file.path(out.dir, 'monthly_death.png') )

```

Apply the code above and have the plot Fig 1.

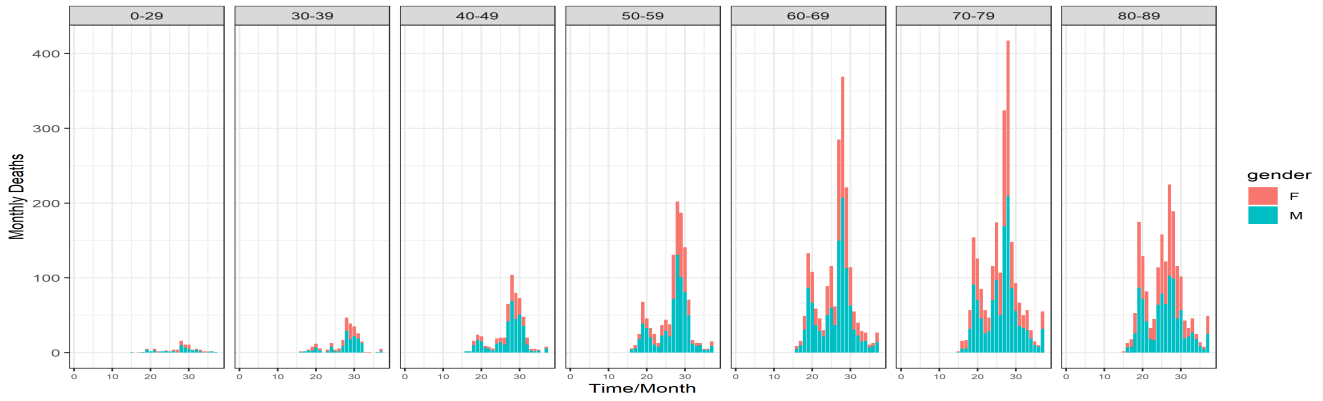


Figure 1: Plot the monthly, reported COVID-19 deaths in the city that you are assigned

3.4 Calculate the cumulating, monthly reported COVID-19 deaths in the city that you are assigned, again by gender and age group.

```

1 setkey(mycity_agg, 'month_id')
2 mycity_cumsum <- mycity_agg[,.(Cum_sum_Death = cumsum(Deaths),
3                               month.id = month_id), by = .(age_group, gender)]
4
5 p <- ggplot(mycity_cumsum, aes(x = month.id, y = Cum_sum_Death, colour = gender)) +
6   geom_line() + facet_grid(~age_group) +
7   theme_bw() + labs(x = 'Time/Month', y = 'Monthly Deaths') +
8   theme(axis.text.x = element_text(size = 6))
9 ggsave(file.path(out.dir, 'cumsum.png'), p, w = 5, h = 5)

```

Fig 2 shows the calculation of cumulating, monthly reported COVID-19 deaths in Belo Horizonte and the table below is the data from the dataset *mycity_cumsum*.

	gender	age_group	month_id	Deaths
1	M	80-89	33.00	663
2	M	80-89	34.00	405
3	M	80-89	35.00	164
4	M	80-89	36.00	104
5	M	80-89	37.00	432

Table 2: The last 5 elements in the dataset *mycity_cumsum*

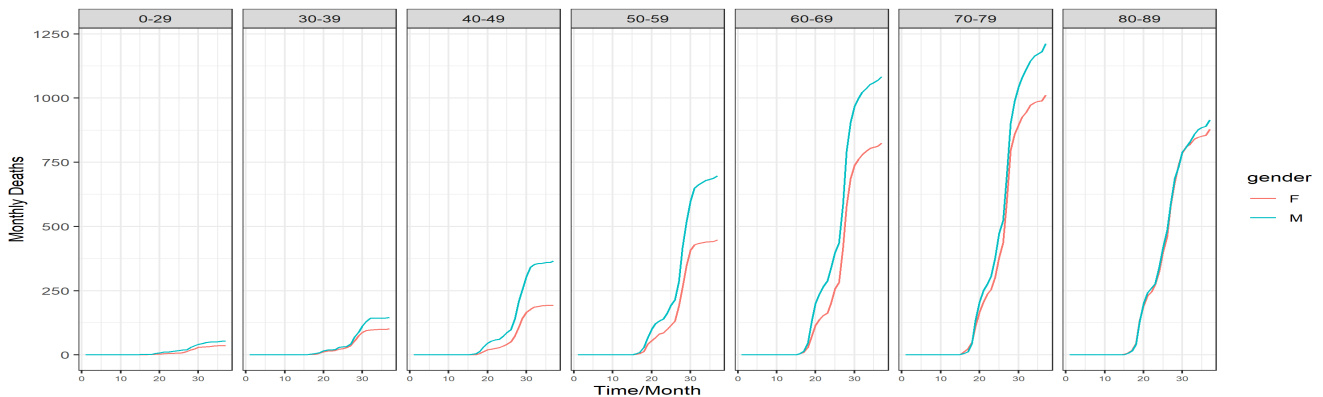


Figure 2: Cumulating monthly reported COVID-19 Deaths in Belo Horizonte

3.5 Plot for the cumulating, monthly reported COVID-19 deaths and the cumulating monthly excess deaths for each gender and age group in the city

```

1 dfcd_base <- subset(dfcd, year == 2019)
2 dfcd_base <- dfcd_base[, list(deaths_2019 = mean(deaths)), by = c('city', 'month_id', 'month', 'gender', 'age_group')]
3 set(dfcd_base, NULL, 'month_id', NULL)
4 dfcd2 <- merge(dfcd, dfcd_base, by = c('city', 'month', 'gender', 'age_group'))
5 dfcd2[, exc_deaths := deaths - deaths_2019]
6 mycity2<- dfcd2[which(city == "Belo Horizonte"),]
7
8 mycity2_ttagg <- mycity2[,.(Deaths = sum(deaths),
9                               exDeaths = sum(exc_deaths)), by = .(gender,age_group, month_id)]
10 mycity2_ttcumsum <- mycity2_ttagg[, .(CumDeath = cumsum(Deaths),
11                                     CumExcDeath=cumsum(exDeaths),
12                                     month_id = month_id), by = .(gender,age_group)]
13 mycity2_ttcumsum <- data.table::melt(mycity2_ttcumsum, id.vars = c('gender','age_group','month_id'))
14
15 p <- ggplot(mycity2_ttcumsum , aes(x = month_id, y = value, fill =variable)) +
16   geom_col() + facet_grid(gender~age_group) +
17   theme_bw()+
18   labs(
19     x='Month Index',
20     y='Numer of Death')
21
22 ggsave(file.path( out.dir, 'CumDexD.png'),p, w=5, h=5)

```

Fig 3 is as required.

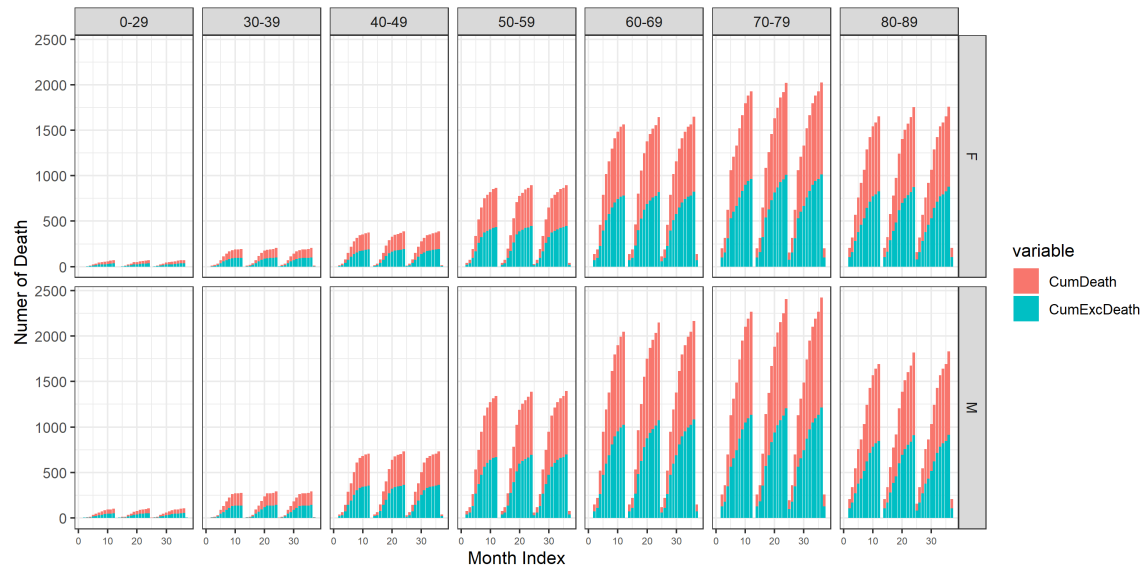


Figure 3: Cumulating and Culmulating Excess Deaths monthly reported in Belo Horizonte

3.6 Calculate the cumulating, monthly reported COVID-19 deaths in the city that you are assigned, per 1,000 population in each gender and age group.

```

1 # load population size projections.
2 data <- as.data.table(read.csv('PNADc_2020.csv',encoding = 'latin1'))
3 age_group_data <- c('0-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80-89')

```

```

4 cut_breaks <- c(-0.1,29,39,49,59,69,79,89)
5 data <-subset(data, age <90)
6 data$age <- cut(x = data$age, breaks = cut_breaks,
7               labels = age_group_data,
8               right = TRUE)
9
10 data <- data[!is.na(data$age),]
11 data <- data[, .(population = sum(population)),by = c('city','sex','age')]
12 setnames(data,'age', 'age_group')
13 setnames(data,'sex', 'gender')
14 set(data,which(data$gender == 'Women'), 'gender', 'F')
15 set(data,which(data$gender == 'Men'), 'gender', 'M')
16 mycity_data <- data[which(city =='Belo Horizonte'),]
17 mycity_cumsumpp <- merge(mycity_cumsum, mycity_data , by= c('gender','age_group'))
18 set(mycity_cumsumpp,NULL,'city', NULL)
19 deathperthousand <- mycity_cumsumpp[, Ppdeath :=1000*Cum_sum_Death/population]

```

deathperthousand stores the wanted data:

	gender	age-group	Cum_sum_Death	month.id	population	Ppdeath
1	M	80-89	687	28.00	24293	28.28
2	M	80-89	733	29.00	24293	30.17
3	M	80-89	790	30.00	24293	32.52
4	M	80-89	809	31.00	24293	33.30
5	M	80-89	831	32.00	24293	34.21

Table 3: The last 10 data in *deathperthousand*

3.7 Make a suitable plot that shows the total reported COVID-19 deaths by the end of 2021 across all cities, per 1,000 population in each city. Show cities on the x-axis and the total per 1,000 population on the y-axis.

```

1 population_total <-data[, .(Population = sum(population)), by = .(city)] # 27 rows
2
3 dfcd_21 <- dfcd[which(year == 2021),], .(Death =sum(deaths)), by = .(city)]
4 dfcd_21pp <- merge(dfcd_21, population_total, by = c('city'))
5 dfcd_21agg<- dfcd_21pp[, DeathperT := 1000*sum(Death)/Population, by = .(city)]
6
7 p <- ggplot(dfcd_21agg, aes(x = city, y = DeathperT)) + geom_col()+
8   theme_bw() + theme(axis.text.x = element_text(angle = 60,size = 6)) +
9     labs( x='City', y='Death Per one Thousand Population') +
10     theme(axis.text.x =element_text(size = 6))
11
12 ggsave(file.path(out.dir, 'dth_per_thousand.png'),w = 10, h = 5)
13 include_graphics(file.path(out.dir, 'dth_per_thousand.png'))

```

Here we use the inner map to map the two dataset according to the name of the city. But since each computer decrypts different for the name of the city (i.e. city name Sao Luis loaded from *civil_registry_covid_cities_detailed.csv* is shown "S*o Lu*s" on my computer but from the dataset PNADc_2020 is "São Paulo"), there are a few cities not matched.

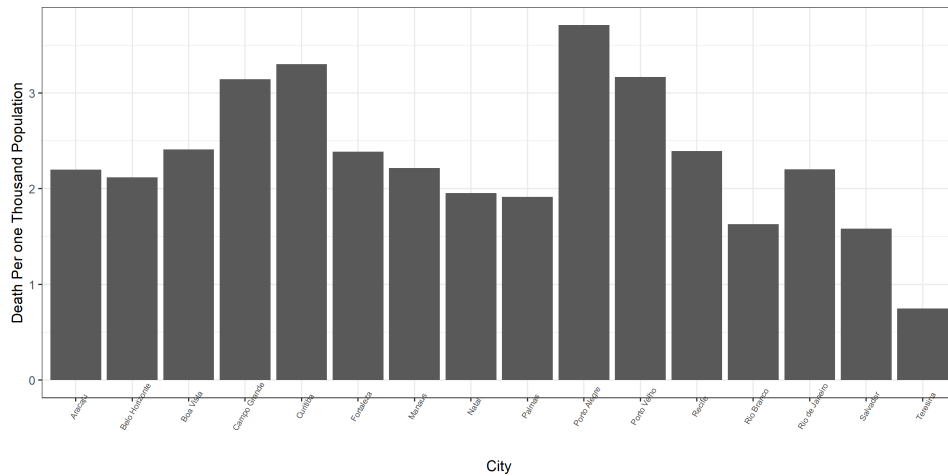


Figure 4: The total reported COVID-19 deaths by the end of 2021 across all cities, per 1,000 population in each city

3.8 Comment if this plot is a fair comparison of total reported COVID-19 deaths across cities.

No. Because the total amount of population for each city is different. For small population, if there are a few people dead, then it would tend to have a bigger proportion compared with the same amount of death in a larger population. So the city with smaller population tends to have a larger proportion of deaths.

4 Assessing the Coursework 1 model (6 points)

First we construct the dataset *stan_data* and generate the posterior estimates and save the data into *m1_fit*.

```

1 file <- file.path(data.dir, 'civil_registry_covid_cities_detailed.csv')
2 dd <- as.data.table(read.csv(file))
3 dd <- subset(dd, select = -c(state, state_ibge_code, city_ibge_code, created_at))
4 dd <- melt(dd, id.vars = c('date', 'city', 'place', 'gender', 'age_group'))
5 set(dd, dd[, which(is.na(value))], 'value', 0L)
6
7 # summarize daily deaths by city gender age
8 dd <- dd[, list(deaths = sum(value)), by = c('date', 'city', 'gender', 'age_group')]
9
10 # define date as Date object
11 set(dd, NULL, 'date', dd[, as.Date(date)])
12
13 # define months deaths, and month index
14 dw <- data.table(date = dd[, seq(min(date), max(date), 1)])
15 dw[, month := as.integer(strftime(date, format = '%m'))]
16 dw[, year := as.integer(strftime(date, format = '%Y'))]
17 dw[, month_id := (year - 2019) * 12 + month]
18 tmp <- dw[, list(month_start = min(date)), by = 'month_id']
19 dw <- merge(dw, tmp, by = 'month_id')
20 dd <- merge(dd, subset(dw, select = -c(month, year)), by = 'date')
21
22 # specify age groups as requested
23 set(dd, dd[, which(age_group == '9-')], 'age_group', '0-29')
24 set(dd, dd[, which(age_group == '10-19')], 'age_group', '0-29')
25 set(dd, dd[, which(age_group == '20-29')], 'age_group', '0-29')
26 set(dd, dd[, which(age_group == '90-99')], 'age_group', '>=90')
27 set(dd, dd[, which(age_group == '100+')], 'age_group', '>=90')

```

```

28 set(dd, dd[, which(is.na(age_group))], 'age_group', 'unknown')
29 dd <- subset(dd, ! age_group %in% c('>=90','unknown'))
30
31 # define gender as factor
32 set(dd, NULL, 'gender', dd[, factor(gender, levels = c('M','F'), labels = c('Men','Women'))])
33
34 # count deaths by month gender age
35 dd <- dd[, list(deaths = sum(deaths)), by = c('month_id','month_start','city','gender','age_group')]
36
37 # define age as factor
38 tmp <- dd[,sort(unique(age_group))]
39 set(dd, NULL, 'age_group', dd[, factor(age_group, levels = tmp)])
40
41 mycity <- dd[which(city == 'Belo Horizonte'),]
42
43 data <- as.data.table(read.csv('PNADc_2020.csv',encoding = 'latin1'))
44 age_group_data <- c('0-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80-89')
45 cut_breaks <- c(-0.1,29,39,49,59,69,79,89)
46 data <-subset(data, age <90)
47
48 data$age <- cut(x = data$age, breaks = cut_breaks,
49               labels = age_group_data,
50               right = TRUE,)
51
52 data <- data[!is.na(data$age),]
53 data <- data[, .(population = sum(population)),by = c('city','sex','age')]
54 setnames(data,'age', 'age_group')
55 setnames(data,'sex','gender')
56 mycity_data <- data[which(city == 'Belo Horizonte'),]
57 dpop <- data
58 #data populaiton & deaths
59 dds <- merge(mycity,dpop, by = c("city","age_group", 'gender'))
60
61 ddsF <- dds[which(gender == 'Women'),]
62 tmp <- unique(subset(ddsF, select = age_group))
63 setkey(tmp, age_group)
64 tmp[, age_group_id := 1:nrow(tmp)]
65 ddsF <- merge(ddsF, tmp, by = 'age_group')
66 ddsF[, id := 1:nrow(ddsF)]
67 tmp <- subset(ddsF, select = c(id, age_group_id))
68 tmp <- dcast.data.table(tmp, id ~ age_group_id, value.var = 'age_group_id')
69 setnames(tmp, 1 + 1:max(ddsF$age_group_id), paste0('AGE',1:max(ddsF$age_group_id)))
70 for (x in paste0('AGE',1:max(ddsF$age_group_id)))
71 {
72   set(tmp, which(!is.na(tmp[[x]])), x, 1L)
73   set(tmp, which(is.na(tmp[[x]])), x, 0L)
74 }
75 tmp <- unname(as.matrix(subset(tmp, select = -c(id,AGE1))))
76
77 # make binary dummy variables for each month except first
78 tmp2 <- subset(ddsF, select = c(id, month_id))
79 tmp2 <- dcast.data.table(tmp2, id ~ month_id, value.var = 'month_id')
80 setnames(tmp2, 1 + 1:max(ddsF$month_id), paste0('MO',1:max(ddsF$month_id)))
81 for (x in paste0('MO',1:max(ddsF$month_id)))
82 {
83   set(tmp2, which(!is.na(tmp2[[x]])), x, 1L)
84   set(tmp2, which(is.na(tmp2[[x]])), x, 0L)
85 }
86 tmp2 <- unname(as.matrix(subset(tmp2, select = -c(id,MO1))))

```



```

87 tmp <- cbind(tmp, tmp2)
88 design_matrix_X <- tmp
89
90 stan_data <- list()
91 stan_data$N <- nrow(ddsF)
92 stan_data$K <- ncol(design_matrix_X )
93 stan_data$y <- ddsF$deaths
94 stan_data$offset <- log(ddsF$population)
95 stan_data$X <- design_matrix_X
96
97 poi_regression_txt <- "
98 data{
99     int<lower=1> N;
100     int<lower=1> K;
101     int<lower=0> y[N];
102     vector[N] offset;
103     matrix[N,K] X;
104 }
105 parameters{
106     real beta0;
107     vector[K] beta;
108 }
109 transformed parameters{
110     vector[N] log_lambda;
111     log_lambda = beta0 + X * beta + offset;
112 }
113 model{
114     beta0 ~ normal( 0 , 10 );
115     beta ~ normal( 0 , 1 );
116     y ~ poisson_log( log_lambda );
117 }
118 "
119 if (!file.exists(file.path(out.dir, paste0('poi_deaths_fit_',city_of_student_ascii,'.rds'))))
120 {
121     poi_regression_compiled <- rstan::stan_model(
122         model_name = 'poi_regression',
123         model_code = gsub('\t',' ',poi_regression_txt)
124     )
125     #m1_init_beta0 <- subset(dds, age_group_id == 1 & month_id == 1)[, log(deaths/population)]
126     m1_fit <- rstan::sampling(poi_regression_compiled,
127         data = stan_data,
128         warmup = 5e2, iter = 1e4, chains = 4,
129         init = list(list(beta0 = 0 , beta = rep(0, stan_data$K)),
130             list(beta0 = 0 , beta = rep(0, stan_data$K)),
131             list(beta0 = 0 , beta = rep(0, stan_data$K)),
132             list(beta0 = 0 , beta = rep(0, stan_data$K))
133         )
134     )
135     saveRDS(m1_fit, file = file.path(out.dir, paste0('poi_deaths_fit_',city_of_student_ascii,'.rds')))
136 }
137 if (file.exists(file.path(out.dir, paste0('poi_deaths_fit_',city_of_student_ascii,'.rds'))))
138 {
139     m1_fit <- readRDS( file.path(out.dir, paste0('poi_deaths_fit_',city_of_student_ascii,'.rds')) )
140 }

```

4.1 Assess convergence, mixing, and make the trace plot corresponding to the model parameter with smallest effective sample size.

We applied 5000 iterations with 500 iterations as warm-up and built 4 chains. From the code below, we see the effective sample size for β s are no less than 1155.1. There are β_{0-42} with \hat{R} for each parameter close to 1. This implies good convergence for the chain, thereby the estimated values for β s are trustworthy.

```

1 print(m1_fit, digits = 3, prob = c(0.025,0.5,0.975))
2 #               mean se_mean   sd      2.5%      50%      97.5% n_eff  Rhat
3 # beta0         -9.854   0.002 0.052    -9.959    -9.854    -9.751  1155  1.001
4 # beta[1]         0.072   0.001 0.048    -0.023     0.072     0.165  8204  1.000
5 # beta[2]         0.924   0.000 0.039     0.848     0.924     0.999  6204  1.000
6 # beta[3]         1.840   0.000 0.034     1.774     1.840     1.904  5422  1.000
7 # beta[4]         2.564   0.000 0.031     2.504     2.564     2.625  4806  1.000
8 # beta[5]         3.500   0.000 0.030     3.440     3.500     3.559  4775  1.000
9 # beta[6]         4.114   0.000 0.030     4.056     4.114     4.172  4665  1.000
10 ...
11
12 po <- summary(m1_fit)$summary
13 print(min(po[, 'n_eff']))
14 #1155.1
15 print(max(po[, 'Rhat']))
16 #1.001915

```

Fig 5 also shows the convergence for the parameter with the smallest effective sample size.

```

1 po <- rstan::extract(m1_fit, inc_warmup = TRUE, permuted = FALSE)
2 bayesplot::color_scheme_set("mix-blue-pink")
3 #trace plot
4 p <- bayesplot::mcmc_trace(po, pars = "beta0", n_warmup = 5e2,
5   facet_args = list(nrow = 1, labeller = label_parsed))

```



Figure 5: The traceplot for beta0.

There are tota 43 betas but we here only consider the first 6 β s as examples. Fig 9 shows that the β_0 is negatively correlated with β_{1-6} but β_{1-6} are all positively correlated with each other.

```
1 po <- rstan::extract(m1_fit, inc_warmup = FALSE, permuted = FALSE )
2 p <- bayesplot::mcmc_pairs(po, regex_pars = "beta", diag_fun = "dens", off_diag_fun = "hex")
3 ggsave(file.path(out.dir, 'm1_pairplot.pdf'), p, w=20, h=20)
```

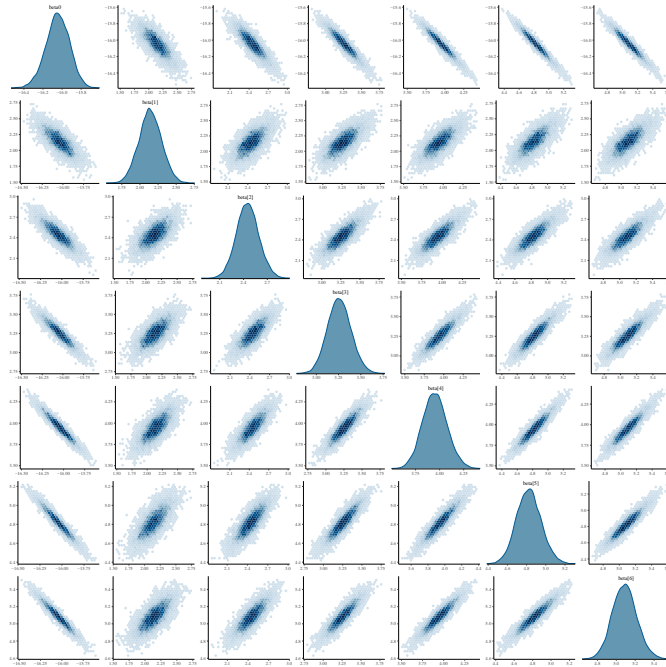


Figure 6: Here we take the pair plot for β_{1-6} . (sub select)

4.2 Make a posterior predictive check and report the proportion of data points inside the 95% posterior predictive credibility intervals associated with each data point, separately for men and women.

```
1 setkey(dds, 'gender')
2 dds[, id := 1:nrow(dds)]
3 ddsFM <- dds
4
5 esti_coef <- rstan::summary(m1_fit)$summary
6 log_lambdaest <- esti_coef[grepl('log_lambda', rownames(esti_coef))][,1]
7 ddsFM <- cbind(ddsFM, log_lambdaest)
8
9 log_poi_fixed_pp <- ddsFM[,
10 {
11   lambda <- exp(log_lambdaest)
12   pred_y <- rpois(10000, lambda)
13   pred_y_stat <- quantile(pred_y, p=c(0.5, 0.025, 0.975), type=1)
14   list(PP_DEATHS_STAT = pred_y_stat, PP_DEATHS_TYPE=c('PP_M', 'PP_CL', 'PP_CU'))
15 },
16 by=c('gender', 'month_id', 'age_group')]
17
18 log_poi_fixed_pp2 <- dcast.data.table(log_poi_fixed_pp, month_id+gender+age_group~PP_DEATHS_TYPE, value.var='PP_DEATHS_STAT')
19
20 # merge observations
```

```

21 tmp2 <- subset(ddsFM, select = c('month_id', 'gender', 'deaths', 'age_group', 'id'))
22 log_poi_fixed_pp3 <- merge(log_poi_fixed_pp2, tmp2, by = c('month_id', 'gender', 'age_group'))
23 set(log_poi_fixed_pp3, which(log_poi_fixed_pp3$gender == 'Women'),
24 'id', log_poi_fixed_pp3[ which(log_poi_fixed_pp3$gender == 'Women'),]$id - 259 + 1)
25
26 p <- ggplot(log_poi_fixed_pp3, aes(x=id)) +
27   geom_point(aes(y=PP_M)) +
28   geom_point(aes(y=deaths), colour='red') +
29   geom_errorbar(aes(ymin=PP_CL, ymax=PP_CU)) +
30   theme_bw() +
31   labs(x='Individuals', y='posterior predictive vs actual deaths\n')
32
33 ggsave(file=file.path(out.dir, 'posterior_pred.png'), p, w=20, h=6)
34 include_graphics(file.path(out.dir, 'posterior_pred.png'))
35
36 ds
37 # posterior predictive check
38 tmp3 <- log_poi_fixed_pp3[which(gender == 'Women'),][,length(which(deaths <= PP_CU & deaths >= PP_CL)) / length(deaths)]
39 tmp3
40 #> 0.9266409
41
42 tmp4 <- log_poi_fixed_pp3[which(gender == 'Men'),][,length(which(deaths <= PP_CU & deaths >= PP_CL)) / length(deaths)]
43 tmp4
44
45 #0.9111969

```

Since we are considering the model

$$\begin{aligned}
Y_{a,t} &\sim \text{Poisson}(\lambda_{a,t}) \\
\log \lambda &= \beta_0 + X\beta + \log P \\
\beta_0 &\sim \mathcal{N}(0, 10^2) \\
\beta &= (\beta_1, \dots, \beta_{A+T-2}) \sim \mathcal{N}(0, 1^2)
\end{aligned}$$

, the first step in a posterior predictive check is to calculate the posterior predictive distribution

$$f(y_j^* | y) = \int l(y_j^* | \theta) f(\theta | y) d\theta$$

, where $\theta = \lambda = \exp(\log \lambda)$ and then we find the probability that an observed data point lies in the 95% credible interval of the corresponding posterior predictive distribution,

$$\frac{1}{n} \sum_j 1(y_j \in [q_{j,\alpha/2}^*, q_{j,1-\alpha/2}^*])$$

where $q_{j,\alpha}^*$ denotes the α quantile of the posterior predictive distribution $f(y_j^* | y)$.

By posterior check, the model for the women has about 92.3% observed data points in the 95% credible interval of the corresponding posterior predictive distribution. Then we re-run the code above again but change the dataset from women to men (change the code `ddsF <- dds[which(gender == 'Women'),]` as `ddsF <- dds[which(gender == 'Men'),]`) and then run the rest of the chunk). For men, it has about 91.1% observed data in the 95% credible interval.

4.3 Comment if the model from Coursework 1 provides a sufficiently good fit to the data for the city that you are assigned.

Referring to Fig 7 and the proportion of data points inside the 95% posterior predictive credibility intervals, we see that the Bayesian model applied to women and men both had high probability coverage for the posterior credible interval so this model fits the data well.

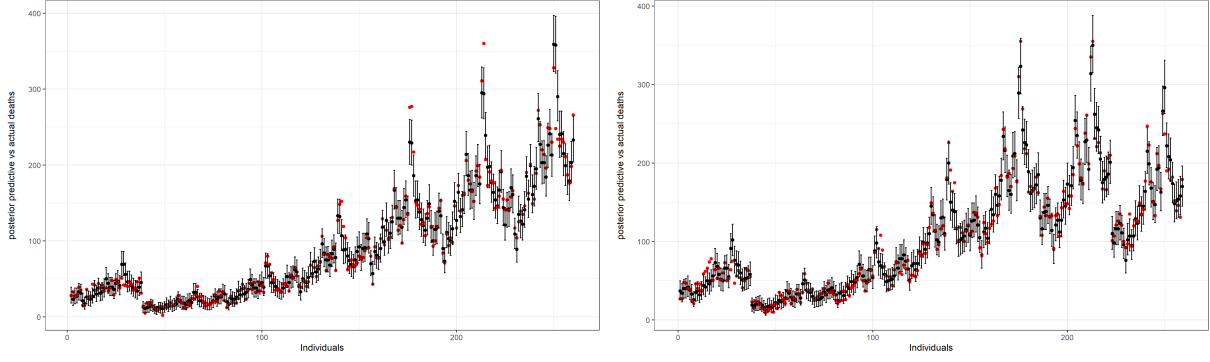


Figure 7: The posterior predictive check for men and women separately in my city (Men right, Women Left). The x-axis is the individual index for men and women separately and y-axis represents the death.

5 Poisson model with random time effects

5.1 Describe Bayesian Poisson model

$$\begin{aligned}
 Y_{a,g,t} &\sim \text{Poisson}(\lambda_{a,g,t}) \\
 \log \lambda_{a,g,t} &= \beta_0 + \beta_1 \text{ age} + \beta_2 \text{ gender} + \beta_{3,a,g} \text{ Time} + \log P_{a,g} \\
 \beta_0 &\sim N(0, 1^2) \\
 \beta_1 &\sim N(0, 1^2) \\
 \beta_2 &\sim N(0, 1^2) \\
 \beta_{3,a,g} &\sim N(\bar{\beta}_{3,a,g}, \sigma_T^2) \\
 \sigma_T &\sim \text{Half-Cauchy}(0, 1) \\
 \bar{\beta}_{3,a,g} &\sim \text{Normal}(0, 11^2)
 \end{aligned}$$

Here we define $Y_{a,g,t}$ to be the number of monthly age- and gender-specific all-cause deaths in my city Belo Horizonte, β_0 is the fixed baseline coefficients for death, β_1 is the fixed effect coefficient for age, β_2 is the fixed effect coefficient for gender, $\beta_{3,a,g}$ is a random effect on Time for the corresponding age and gender, and $P_{a,g}$ is the population for the corresponding age and gender (as we assume the population is independent of the time). I choose the hyperparameter $\bar{\beta}_3$ with standard deviation 11 because there are total 37 months and the `sd` for 1 : 37 is 11. In addition, *gender* takes value $\{0, 1\}$ where 0 represents men and 1 represents women, and *Time* is monthly-recorded with *month.id* (which we no longer treat as factor but real continuous values). I created a matrix for *age*: say there are $A = 7$ age groups, and we allocate only one fixed effect to the age groups $a = 2, \dots, A$, which is similar to the design matrix X but without the month column.

5.2 Implement the hierarchical Bayesian Poisson model in Stan.

```

1 #Data
2 dpop <- data
3 dds <- merge(mycity, dpop, by = c("city", "age_group", "gender"))
4 tmp <- unique(subset(dds, select = age_group))
5 setkey(tmp, age_group)
6 tmp[, age_group_id := 1:nrow(tmp)]
7 dds <- merge(dds, tmp, by = "age_group")
8 tmp <- unique(subset(dds, select = gender))
9 setkey(tmp, gender)
10 tmp[, gender_id := 0:(nrow(tmp)-1)]
11 dds <- merge(dds, tmp, by = "gender")
12 dds[, id := 1:nrow(dds)]
13 ageM <- model.matrix(~ -1 + age_group, data = dds)
14 ageM <- design_matrix_X[, -1]
15 # separate random effects using Stan syntax
16 log_poi_RD_txt <- "

```

```

17 data{
18   int<lower=1> N; // number of observations
19   int<lower=1> K; // number of units
20   int<lower=0> y[N];
21   vector[N] age; //
22   vector[N] gender;
23   vector[N] time;
24   int<lower=1> Nk_max; // max number of observations for a unit
25   int<lower=1, upper=Nk_max> units_to_obs_length[K]; // number of observations per unit
26   int<lower=0, upper=N> units_to_obs[K, Nk_max]; //index of observations per unit, with the rest set to 0
27 }
28 parameters{
29
30   real beta0; // baseline
31   real beta1; // age
32   real beta2; // gender
33   // random effects on treatment effect
34   real beta3[K];
35   real<lower=0> beta3_hyper_sd;
36   real beta3_hyper_mean;
37 }
38 transformed parameters{
39   vector[N] obs_log_lambda;
40
41   // fixed effects
42   obs_log_lambda= beta0 + beta1*age + beta2 * gender + beta3_hyper_mean * time ;
43   //
44   // add random effects
45   for(k in 1:K)
46   {
47     //
48     // add baseline random effects and treatment effects
49     obs_log_lambda[ units_to_obs[k, 1:units_to_obs_length[k] ] ] +=
50     (
51       time[ units_to_obs[k, 1:units_to_obs_length[k] ] ] * beta3[k]
52     );
53   }
54 }
55 model{
56   beta3_hyper_mean ~ normal(0,11);
57   beta3_hyper_sd ~ cauchy(0,1);
58   beta3 ~ normal( 0, beta3_hyper_sd);
59   beta0 ~ normal(0,1);
60   beta1 ~ normal(0,1);
61   beta2 ~ normal(0,1);
62   y ~ poisson_log(obs_log_lambda);
63 }
64 "
65
66
67 # define data in format needed for model specification
68
69 stan_data <- list()
70 stan_data$N <- nrow(dds)
71 stan_data$K <- max(dds$month_id)
72 stan_data$y <- dds$deaths
73 stan_data$age <- age
74 stan_data$time <- dds$month_id
75 stan_data$gender <- dds$gender_id

```

```

76 stan_data$Nk_max <- length(unique(dds$age_group_id))*2
77 # define number of observations per unit
78 tmp <- dds[, list(LEN=length(age_group_id)), by='month_id']
79 tmp <- tmp[order(month_id)]
80 stan_data$units_to_obs_length <- tmp$LEN
81 # define index of observations per unit, with the rest set to 0
82 tmp <- matrix(data=0, nrow=stan_data$K, ncol=stan_data$Nk_max)
83 for(k in 1:stan_data$K)
84 {
85   tmp2 <- dds[,which(month_id==k)]
86   tmp[k, 1:length(tmp2)] <- tmp2
87   stopifnot(length(tmp2)==stan_data$units_to_obs_length[k])
88 }
89 stan_data$units_to_obs <- tmp
90
91 # compile the model
92 log_poi_RD_compiled <- rstan::stan_model(
93   model_name= 'log_poi_RD_model',
94   model_code = gsub('\t', ' ', log_poi_RD_txt)
95 )
96
97 # run Stan
98 log_poi_RD_fit <- rstan::sampling(log_poi_RD_compiled,
99   data=stan_data,
100   warmup=1e3, iter=1e4, chains=4
101 )
102
103 # save
104 saveRDS(log_poi_RD_fit, file=file.path(out.dir, 'log_poi_RD_fit.rds'))
105
106 po <- rstan::summary(log_poi_RD_fit)
107 print(po)

```

5.3 Check convergence and mixing

We applied 5000 iterations with 1000 iterations as warm-up and built 4 chains. From the code below, we see the smallest effective sample size is 1042. There are $\beta_0, \beta_1[1 : 7]$, β_2 , $\beta_3[1 : 37]$ and hyperparameter $\sigma, \bar{\beta}_3$ with \hat{R} for each parameter close to 1. This implies good convergence for the chain, thereby the estimated values for $\hat{\beta}$ s are trustworthy.

```

1 print(log_poi_RD_fit, digits = 3, prob = c(0.025,0.5,0.975))
2 # Inference for Stan model: log_poi_Random.
3 # 4 chains, each with iter=10000; warmup=1000; thin=1;
4 # post-warmup draws per chain=4000, total post-warmup draws=16000.
5 #
6 #               mean se_mean   sd      2.5%      50%      97.5% n_eff Rhat
7 # beta0          2.840    0.000 0.030     2.783     2.840     2.900  4375 1.001
8 # beta1[1]        0.308    0.000 0.002     0.303     0.308     0.313 1482 1.003
9 # ...
10 # beta2          -0.146    0.000 0.009    -0.163    -0.146    -0.129 1042 1.004
11 # beta3[1]        0.010    0.000 0.016    -0.020     0.010     0.042 15135 1.000
12 # beta3[2]       -0.020    0.000 0.013    -0.048    -0.020     0.005 12853 1.000
13 # beta3[3]        0.013    0.000 0.010    -0.007     0.013     0.032 13284 1.000
14 # beta3[4]        0.009    0.000 0.008    -0.007     0.009     0.025 13648 1.000
15 # beta3[5]        0.025    0.000 0.007     0.012     0.025     0.038 13087 1.000
16 # beta3[6]        0.007    0.000 0.006    -0.005     0.007     0.018 11602 1.000
17 # beta3[7]       -0.030    0.000 0.006    -0.041    -0.030    -0.019 11542 1.000

```

```

18 # ...
19 # beta3_hyper_sd      0.016  0.000 0.002    0.013    0.016    0.022 12172 1.000
20 # beta3_hyper_mean    0.017  0.000 0.004    0.010    0.017    0.025 3860 1.002
21 #...
22 po <- summary(log_poi_RD_fit)$summary
23 est <- po[grepl("hyper|beta",rownames(po_Sp)),]
24 min(est[, 'n_eff'])
25 # 1042
26 max(est[, 'Rhat'])

```

Take $\beta_0, \beta_1[1]$ for example. By Fig 8, the graph shows good convergence

```

1 po <- rstan::extract(log_poi_RD_fit, inc_warmup = TRUE, permuted = FALSE)
2 bayesplot::color_scheme_set("mix-blue-pink")
3 #trace plot
4 p <- bayesplot::mcmc_trace(po, pars = c("beta0", "beta1[1]"), n_warmup = 1e3, facet_args = list(nrow = 2, labeller = label_parsed))
5 ggsave(file.path(out.dir, 'RDtraceplot.png'), p, w = 10, h = 5)
6 include_graphics(file.path(out.dir, 'RDtraceplot.png'))

```

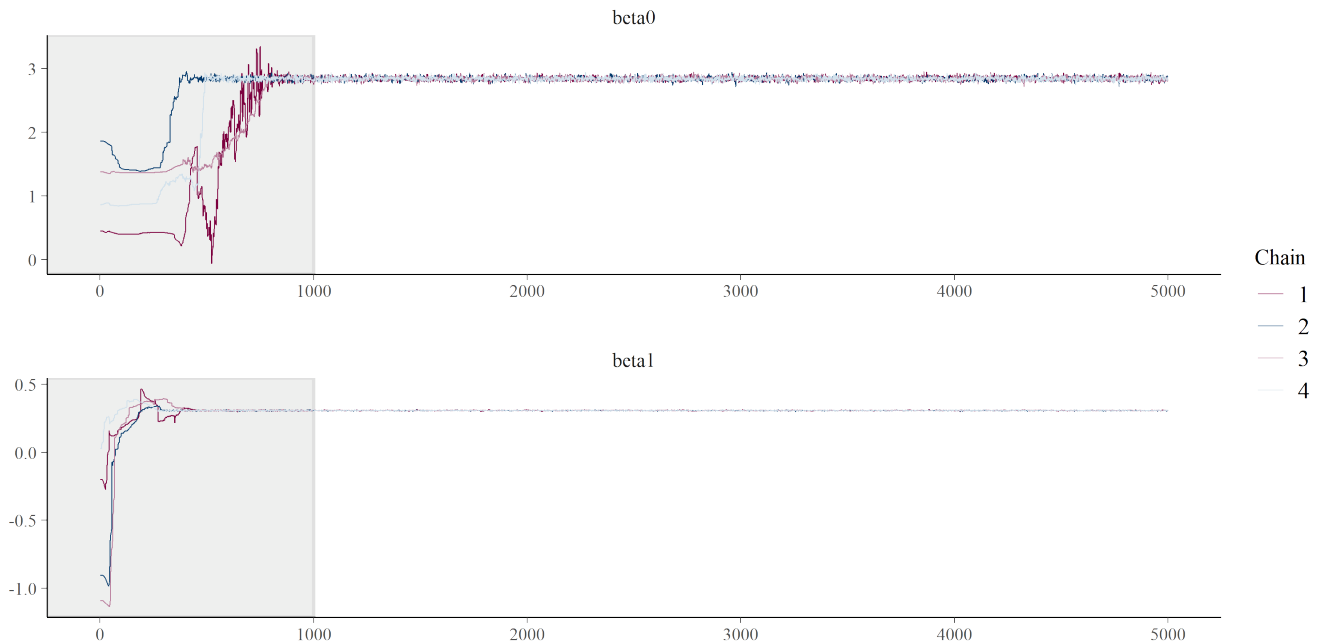


Figure 8: The traceplot for $\beta_0, \beta_1[1]$.

5.4 Make a posterior predictive check and report the proportion of data points inside the 95% posterior predictive credibility intervals associated with each data point, separately for men and women.

```

1 ddsFM <- dds
2 esti_coef <- rstan::summary(log_poi_RD_fit)$summary
3 log_lambdaest <- esti_coef[grepl('log_lambda', rownames(esti_coef))][,1]
4 ddsFM <- cbind(ddsFM, log_lambdaest)
5
6 log_poi_fixed_pp <- ddsFM[,
7 {

```



```

8   lambda <- exp(log_lambdaest)
9   pred_y <- rpois(10000,lambda)
10  pred_y_stat <- quantile(pred_y, p=c(0.5,0.025,0.975), type=1)
11  list(PP_DEATHS_STAT = pred_y_stat, PP_DEATHS_TYPE=c('PP_M', 'PP_CL', 'PP_CU'))
12 },
13  by=c('gender', 'month_id', 'age_group')]
14
15 log_poi_fixed_pp2 <- dcast.data.table(log_poi_fixed_pp, month_id+gender+age_group~PP_DEATHS_TYPE, value.var='PP_DEATHS_STAT')
16
17 # merge observations
18 tmp2 <- subset(ddsFM, select = c('month_id', 'gender', 'deaths', 'age_group', 'id'))
19 log_poi_fixed_pp3 <- merge(log_poi_fixed_pp2, tmp2, by = c('month_id', 'gender', 'age_group'))
20 set(log_poi_fixed_pp3, which(log_poi_fixed_pp3$gender == 'Women'),
21    'id', log_poi_fixed_pp3[ which(log_poi_fixed_pp3$gender == 'Women'),]$id - 259 +1)
22
23 # posterior predictive check
24 tmp3 <- log_poi_fixed_pp3[which(gender == 'Women'),][,length(which(deaths <= PP_CU & deaths >= PP_CL)) / length(deaths)]
25 tmp3
26 #> 0.7490347
27
28 tmp4 <- log_poi_fixed_pp3[which(gender == 'Men'),][,length(which(deaths <= PP_CU & deaths >= PP_CL)) / length(deaths)]
29 tmp4
30
31 #0.7837838

```

By posterior predictive check, we see that for women, there is only about 74.9% of data within the 95% posterior predictive credibility intervals and 78.4% for men.

5.5 Plot the posterior predictive checks

```

1  p <- ggplot(log_poi_fixed_pp3, aes(x=month_id)) +
2    geom_point(aes(y=PP_M)) +
3    geom_point(aes(y=deaths), colour='red') +
4    geom_errorbar(aes(ymin=PP_CL, ymax=PP_CU)) +
5    theme_bw() +
6    labs(x='Individuals', y='posterior predictive vs actual deaths\n') +
7    facet_grid(gender~age_group)
8
9  ggsave(file=file.path(out.dir, 'posterior_pred_RD.png'), p, w=20, h=6)
10 include_graphics(file.path(out.dir, 'posterior_pred_RD.png'))

```

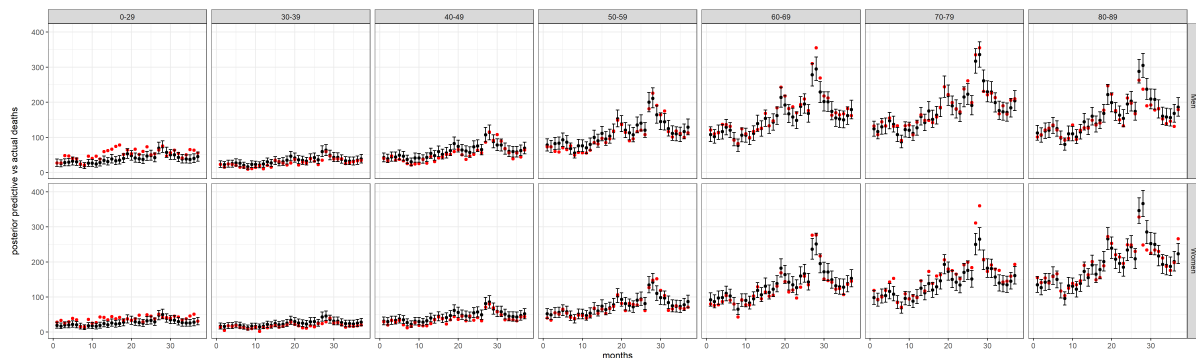


Figure 9: The posterior predictive checks with time on the x-axis, and age and gender in row facets

5.6 Comment

By the posterior check, we see that there are over 25% of observed data outside the 95% posterior predictive credibility intervals and over 21% for men. Even though this model is valid for the city applied, the model is less as good as the model from Coursework1.

6 Your best model

We construct "simplest plausible Bayesian Poisson model with the fewest possible number of variables". Let us denote observed all cause deaths among men in the city that you are given, month t and age group a by $Y_{a,t}$, and denote the total number of age groups by A , such that $a = 1, \dots, A$, and denote the total number of months by T , such that $t = 1, \dots, T$. In total, we have $A * T * 2$ observations. We then model the count data with the Poisson model

$$\begin{aligned} Y_{a,t} &\sim \text{Poisson}(\lambda_{a,g,t}) \\ \log \lambda &= \beta_0 + X\beta + \log P \\ \beta_0 &\sim \mathcal{N}(0, 10^2) \\ \beta &= (\beta_1, \dots, \beta_{A+T-2}) \sim \mathcal{N}(0, 1^2) \end{aligned}$$

where $\lambda_{a,g,t}$ is the expected number of all cause deaths in this city in age group a , gender g and month t ; P is an $A * T * 2$ dimensional column vector of age-specific population sizes in this city, which for deaths in age group a gives the population size in age group a ; β_0 is a real-value baseline regression parameter; and β is a $A + T - 1$ dimensional vector of regression parameter, one for each covariate. The first $A + T - 2$ columns of the design matrix X are defined the same as for coursework 1 but add a column 'gender' which has value 0 for men and 1 for women. And the matrix X is therefore has double rows as in coursework 1.

6.1 Implement your hierarchical Bayesian Poisson model in Stan.

Data Construction::

```
1 file <- file.path(data.dir, 'civil_registry_covid_cities_detailed.csv')
2 dd <- as.data.table(read.csv(file))
3 dd <- subset(dd, select = -c(state, state_ibge_code, city_ibge_code, created_at))
4 dd <- melt(dd, id.vars = c('date', 'city', 'place', 'gender', 'age_group'))
5 set(dd, dd[, which(is.na(value))], 'value', 0L)
6
7 # summarize daily deaths by city gender age
8 dd <- dd[, list(deaths = sum(value)), by = c('date', 'city', 'gender', 'age_group')]
9
10 # define date as Date object
11 set(dd, NULL, 'date', dd[, as.Date(date)])
12
13 # define months deaths, and month index
14 dw <- data.table(date = dd[, seq(min(date), max(date), 1)])
15 dw[, month := as.integer(strftime(date, format = '%m'))]
16 dw[, year := as.integer(strftime(date, format = '%Y'))]
17 dw[, month_id := (year - 2019) * 12 + month]
18 tmp <- dw[, list(month_start = min(date)), by = 'month_id']
19 dw <- merge(dw, tmp, by = 'month_id')
20 dd <- merge(dd, subset(dw, select = -c(month, year)), by = 'date')
21
22 # specify age groups as requested
23 set(dd, dd[, which(age_group == '9-')], 'age_group', '0-29')
24 set(dd, dd[, which(age_group == '10-19')], 'age_group', '0-29')
25 set(dd, dd[, which(age_group == '20-29')], 'age_group', '0-29')
26 set(dd, dd[, which(age_group == '90-99')], 'age_group', '>=90')
27 set(dd, dd[, which(age_group == '100+')], 'age_group', '>=90')
28 set(dd, dd[, which(is.na(age_group))], 'age_group', 'unknown')
```

```

29 dd <- subset(dd, ! age_group %in% c('>=90','unknown'))
30
31 # define gender as factor
32 set(dd, NULL, 'gender', dd[, factor(gender, levels = c('M','F'), labels = c('Men','Women'))])
33
34 # count deaths by month gender age
35 dd <- dd[, list(deaths = sum(deaths)), by = c('month_id','month_start','city','gender','age_group')]
36
37 # define age as factor
38 tmp <- dd[,sort(unique(age_group))]
39 set(dd, NULL, 'age_group', dd[, factor(age_group, levels = tmp)])
40
41 data <- as.data.table(read.csv('PNADc_2020.csv',encoding = 'latin1'))
42 age_group_data <- c('0-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80-89')
43 cut_breaks <- c(-0.1,29,39,49,59,69,79,89)
44 data <-subset(data, age <90)
45
46 data$age <- cut(x = data$age, breaks = cut_breaks,
47               labels = age_group_data,
48               right = TRUE,
49               )
50
51 data <- data[!is.na(data$age),]
52 data <- data[, .(population = sum(population)),by = c('city','sex','age')]
53 setnames(data,'age', 'age_group')
54 setnames(data,'sex','gender')
55 dpop <- data
56 #data populaiton & deaths
57 dds <- merge(dd,dpop, by = c("city","age_group", 'gender'))
58 city_of_student_ascii = 'Belo Horizonte'
59 dds <- dds[which(city == 'Belo Horizonte'),]
60
61 design_matrix_X <- model.matrix(~ -1 + age_group + as.factor(month_id) +as.factor(gender), data = dds)
62 design_matrix_X <- design_matrix_X[,-1]
63
64
65 stan_data <- list()
66 stan_data$N <- nrow(dds)
67 stan_data$K <- ncol(design_matrix_X )
68 stan_data$y <- dds$deaths
69 stan_data$offset <- log(dds$population)
70 stan_data$X <- design_matrix_X

```

Fit into the Rstan model:

```

1 poi_regression2_txt <- "
2 data{
3   int<lower=1> N;
4   int<lower=1> K;
5   int<lower=0> y[N];
6   vector[N] offset;
7   matrix[N,K] X;
8 }
9 parameters{
10   real beta0;
11   vector[K] beta;
12 }
13 transformed parameters{

```

```

14     vector[N] log_lambda;
15     log_lambda = beta0 + X * beta + offset;
16 }
17 model{
18     beta0 ~ normal( 0 , 10 );
19     beta ~ normal( 0 , 1 );
20     y ~ poisson_log( log_lambda );
21 }
22 "
23 if (!file.exists(file.path(out.dir, paste0('poi_deaths_fit2_',city_of_student_ascii,'.rds'))))
24 {
25     poi_regression2_compiled <- rstan::stan_model(
26         model_name = 'poi_regression',
27         model_code = gsub('\t',' ',poi_regression2_txt)
28     )
29     #m1_init_beta0 <- subset(dds, age_group_id == 1 & month_id == 1)[, log(deaths/population)]
30     m2_fit <- rstan::sampling(poi_regression2_compiled,
31         data = stan_data,
32         warmup = 5e2, iter = 1e4, chains = 4,
33         init = list(list(beta0 = 0 , beta = rep(0, stan_data$K)),
34                     list(beta0 = 0 , beta = rep(0, stan_data$K)),
35                     list(beta0 = 0 , beta = rep(0, stan_data$K)),
36                     list(beta0 = 0 , beta = rep(0, stan_data$K))
37                 )
38     )
39     saveRDS(m2_fit, file = file.path(out.dir, paste0('poi_deaths_fit2_',city_of_student_ascii,'.rds')))
40 }
41 if (file.exists(file.path(out.dir, paste0('poi_deaths_fit2_',city_of_student_ascii,'.rds'))))
42 {
43     m2_fit <- readRDS( file.path(out.dir, paste0('poi_deaths_fit2_',city_of_student_ascii,'.rds')) )
44 }

```

6.2 Check convergence and mixing in an appropriate manner

We applied 10000 iterations with 500 iterations as warm-up and built 4 chains. The minimum effective sample size is about 2372.063 and the maximum \hat{R} is 1.001915. There are β_{0-43} and \hat{R} for all parameters are close to 1 which indicates good convergence. Also, Fig 10 shows that the chain converges.

```

1 print(m2_fit)
2 #
3 # beta0          mean se_mean sd      2.5%      25%      50%      75%      97.5% n_eff Rhat
4 # beta[1]        0.15   0.00 0.03      0.10      0.13      0.15      0.18      0.21 14948   1
5 # beta[2]        0.98   0.00 0.02      0.93      0.96      0.98      1.00      1.03 11968   1
6 # beta[3]        1.84   0.00 0.02      1.80      1.83      1.84      1.86      1.89 9864    1
7 # ...
8 # beta[41]       0.37   0.00 0.04      0.28      0.34      0.37      0.39      0.45 3797    1
9 # beta[42]       0.47   0.00 0.04      0.39      0.44      0.47      0.50      0.55 3329    1
10 # beta[43]      -0.43   0.00 0.01     -0.44     -0.43     -0.43     -0.42     -0.41 54284   1
11 po <- summary(m2_fit)$summary
12 print(min(po[, 'n_eff']))
13 #2372.063
14 print(max(po[, 'Rhat']))
15 #1.001915
16
17 po <- rstan::extract(m2_fit, inc_warmup = TRUE, permuted = FALSE)
18 bayesplot::color_scheme_set("mix-blue-pink")
19

```

```

20 p <- bayesplot::mcmc_trace(po, pars = "beta0", n_warmup = 5e2,
21   facet_args = list(nrow = 1, labeller = label_parsed))
22
23 ggsave(file.path(out.dir, "bayesplot2.png"), p, w=10, h=5)
24 include_graphics(file.path(out.dir, "bayesplot2.png"))

```

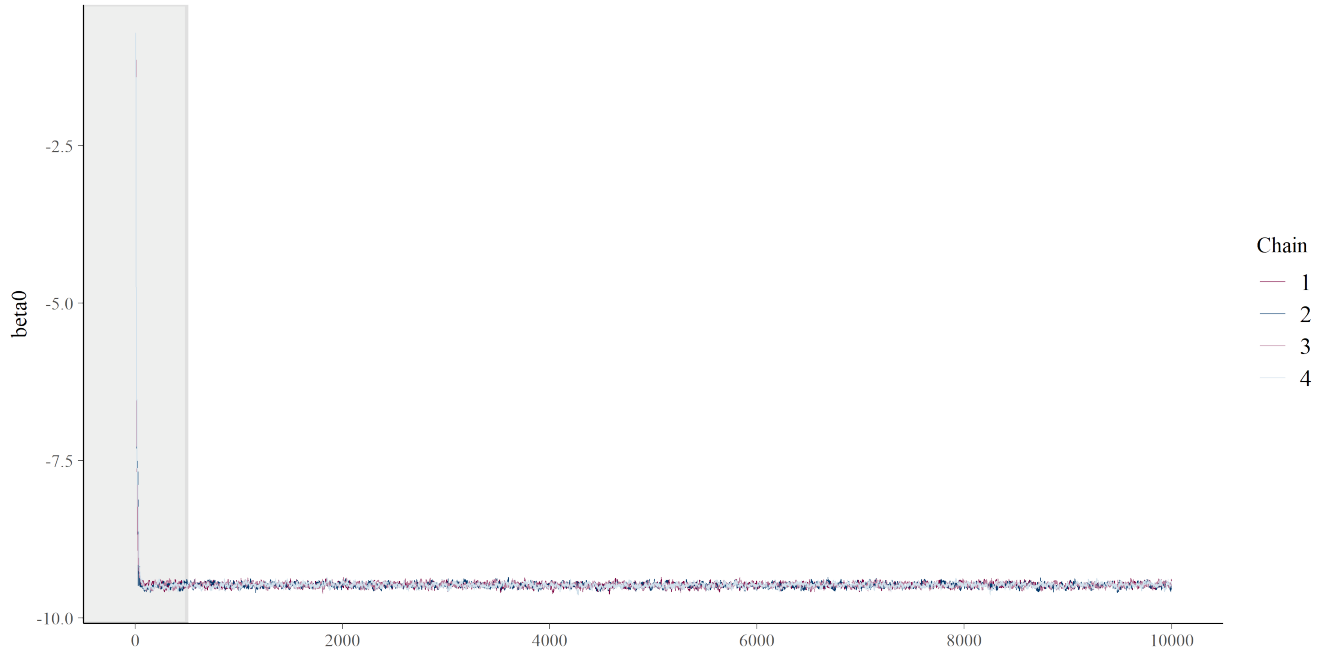


Figure 10: The traceplot for my parsimonious Bayesian Poisson model .

6.3 Quantify model fit using a posterior predictive check.

```

1 dds[, id := 1:nrow(dds)]
2 ddsFM <- dds
3 esti_coef <- rstan::summary(m2_fit)$summary
4 log_lambdaest <- esti_coef[grepl('log_lambda', rownames(esti_coef))][,1]
5 ddsFM <- merge(ddsFM, log_lambdaest)
6 log_poi_fixed_pp <- ddsFM[,
7   {
8     lambda <- exp(log_lambdaest)
9     pred_y <- rpois(10000, lambda)
10    pred_y_stat <- quantile(pred_y, p=c(0.5, 0.025, 0.975), type=1)
11    list(PP_DEATHS_STAT = pred_y_stat, PP_DEATHS_TYPE=c('PP_M', 'PP_CL', 'PP_CU'))
12  },
13   by=c('gender', 'month_id', 'age_group')]
14
15 log_poi_fixed_pp2 <- dcast.data.table(log_poi_fixed_pp, month_id+gender+age_group~PP_DEATHS_TYPE, value.var='PP_DEATHS_STAT')
16
17 # merge observations
18 tmp2 <- subset(ddsFM, select = c('month_id', 'gender', 'deaths', 'age_group', 'id'))
19 log_poi_fixed_pp3 <- merge(log_poi_fixed_pp2, tmp2, by = c('month_id', 'gender', 'age_group'))
20 set(log_poi_fixed_pp3, which(log_poi_fixed_pp3$gender == 'Women'),
21   'id', log_poi_fixed_pp3[ which(log_poi_fixed_pp3$gender == 'Women'),]$id - 259 + 1)
22
23 p <- ggplot(log_poi_fixed_pp3, aes(x=month_id)) +

```

```

24     geom_point(aes(y=PP_M)) +
25     geom_point(aes(y=deaths), colour='red') +
26     geom_errorbar(aes(ymin=PP_CL, ymax=PP_CU)) +
27     theme_bw() +
28     labs(x='Individuals', y='posterior predictive vs actual deaths\n') +
29     facet_grid(gender~age_group)
30 ggsave(file=file.path(out.dir, 'posterior_pred2.png'), p, w=8, h=6)
31 include_graphics(file.path(out.dir, 'posterior_pred2.png'))
32
33 # posterior predictive check
34 tmp3 <- log_poi_fixed_pp3[,length(which(deaths <= PP_CU & deaths >= PP_CL)) / length(deaths)]
35 tmp3
36 #> 0.9034749

```

By posterior predictive check, we see that there are about 90% of the data in the 95% credible interval for the entire data. Fig 11 shows how the credible interval covers the real data points.

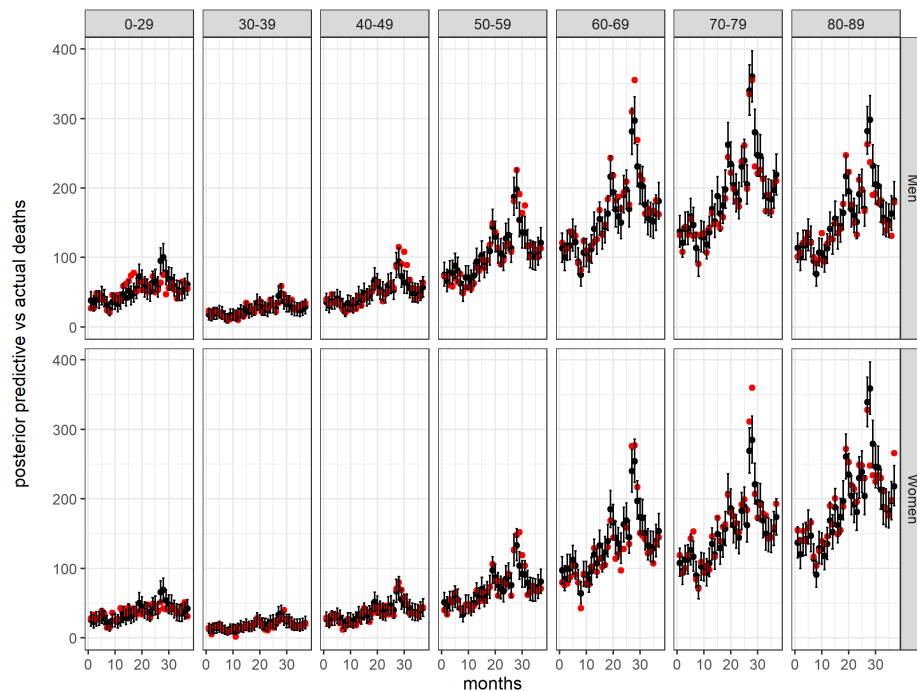


Figure 11: The posterior predictive check plot with month on the x-axis, deaths on the y-axis, facet with gender in row and age group in column.

7 Under-reporting of COVID-19 deaths

For this entire question, we will use the dataset *dfcd* which contains the Covid-19 death, *dd* which contains the death caused by all reasons, and *data* which contains the population for each city, as presented before.

Data For this questions:

```

1 # Total Death
2 ```{r}
3 file <- file.path(data.dir, 'civil_registry_covid_cities_detailed.csv')
4 dd <- as.data.table(read.csv(file))
5 dd <- subset(dd, select = -c(state, state_ibge_code, city_ibge_code, created_at))
6 dd <- melt(dd, id.vars = c('date', 'city', 'place', 'gender', 'age_group'))
7 set(dd, dd[,which(is.na(value))], 'value', 0L)

```

```

8
9 # summarize daily deaths by city gender age
10 dd <- dd[, list(deaths = sum(value)), by = c('date','city','gender','age_group')]
11
12 # define date as Date object
13 set(dd, NULL, 'date', dd[, as.Date(date)])
14
15 # define months deaths, and month index
16 dw <- data.table(date = dd[, seq(min(date),max(date),1)])
17 dw[, month := as.integer(strftime(date, format = '%m'))]
18 dw[, year := as.integer(strftime(date, format = '%Y'))]
19 dw[, month_id := (year - 2019) * 12 + month]
20 tmp <- dw[, list(month_start = min(date)), by = 'month_id']
21 dw <- merge(dw, tmp, by = 'month_id')
22 dd <- merge(dd, subset(dw, select = -c(month,year)), by = 'date')
23
24 # specify age groups as requested
25 set(dd, dd[, which(age_group == '9-')], 'age_group', '0-29')
26 set(dd, dd[, which(age_group == '10-19')], 'age_group', '0-29')
27 set(dd, dd[, which(age_group == '20-29')], 'age_group', '0-29')
28 set(dd, dd[, which(age_group == '90-99')], 'age_group', '>=90')
29 set(dd, dd[, which(age_group == '100+')], 'age_group', '>=90')
30 set(dd, dd[, which(is.na(age_group))], 'age_group', 'unknown')
31 dd <- subset(dd, ! age_group %in% c('>=90','unknown'))
32
33 # define gender as factor
34 set(dd, NULL, 'gender', dd[, factor(gender, levels = c('M','F'), labels = c('Men','Women'))])
35
36 # count deaths by month gender age
37 dd <- dd[, list(deaths = sum(deaths)), by = c('month_id','month_start','city','gender','age_group')]
38
39 # define age as factor
40 tmp <- dd[,sort(unique(age_group))]
41 set(dd, NULL, 'age_group', dd[, factor(age_group, levels = tmp)])
42 ```
43
44 #Population Data
45 ```{r}
46 data <- as.data.table(read.csv('PNADc_2020.csv',encoding = 'latin1'))
47 age_group_data <- c('0-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80-89')
48 cut_breaks <- c(-0.1,29,39,49,59,69,79,89)
49 data <-subset(data, age <90)
50
51 data$age <- cut(x = data$age, breaks = cut_breaks,
52               labels = age_group_data,
53               right = TRUE,
54               )
55
56 data <- data[!is.na(data$age),]
57 data <- data[, .(population = sum(population)),by = c('city','sex','age')]
58 setnames(data,'age', 'age_group')
59 setnames(data,'sex','gender')
60 ```

```

7.1 Calculate the monthly excess deaths for each gender and age group

```

1 dd[, year:= as.integer(substr(month_start, 1,4))]
2 dd[, month:= as.integer(substr(month_start,6,7))]
3 dd_base <- subset(dd, year == 2019)
4 dd_base <- dd_base[, list(deaths_2019 = mean(deaths)), by = c('city','month_id','month','gender', 'age_group')]
5 set(dd_base, NULL, 'month_id', NULL)
6 dd <- merge(dd, dd_base, by = c('city','month','gender', 'age_group'))
7 dd[, exc_deaths := deaths - deaths_2019]
8 mycity <- dd[which(city == 'Belo Horizonte'),]
9 dropcol <- c('city')
10 mycity <- mycity[!dropcol, with = FALSE]
11 xtable(head(subset(dd, select = c('month_id','gender','age_group', 'exc_deaths')),10))

```

mycity contains the wanted data. Here we present the first 10 rows of the calculated monthly excess deaths for each gender and age group as follow:

	month_id	gender	age_group	exc_deaths
1	1.00	Men	0-29	0.00
2	13.00	Men	0-29	-6.00
3	25.00	Men	0-29	-2.00
4	37.00	Men	0-29	-4.00
5	1.00	Men	30-39	0.00
6	13.00	Men	30-39	1.00
7	25.00	Men	30-39	1.00
8	37.00	Men	30-39	3.00
9	1.00	Men	40-49	0.00
10	13.00	Men	40-49	-1.00

7.2 Plot

Here we use the posterior estimated data *log_poi_fixed_pp3* from Q6 and make the plot Fig 12:

```

1 covidExc <-dfcd
2 setnames(covidExc, 'deaths','CovidDeaths')
3 set(covidExc, NULL, 'year',NULL)
4 ddT <- merge(covidExc,dd, by =c('city', 'gender','age_group','month_id'))
5 ddT <- merge(ddT, data, by =c('city', 'gender','age_group'))
6 Mycity <- ddT[which(city == 'Belo Horizonte'),]
7 Mycity[, ExcDeathPer := 1000 * exc_deaths/population]
8 # apply the best model
9 fit_pp <- log_poi_fixed_pp3
10 Mycity2 <- merge(fit_pp, Mycity, by = c('gender','age_group','month_id','deaths'))
11 Mycity2[, UpExcDeathPer := 1000 *(PP_CU- deaths_2019)/population]
12 Mycity2[, LwExcDeathPer := 1000 *(PP_CL- deaths_2019)/population]
13 Mycity2[, MExcDeathPer := 1000 *(PP_M- deaths_2019)/population]
14 MycityCum <- Mycity2[, .(CovidCumDeath= cumsum(CovidDeaths),
15                        CumExcDeathPer = cumsum(ExcDeathPer),
16                        CumUpExcDeathPer = cumsum(UpExcDeathPer),
17                        CumLwExcDeathPer = cumsum(LwExcDeathPer),
18                        CumMExcDeathPer = cumsum(MExcDeathPer),
19                        month_id = month_id), by = c('gender','age_group')]
20 p <- ggplot(MycityCum, aes(x =CovidCumDeath)) +
21   geom_point(aes(y=CumExcDeathPer, colour=gender)) +
22   geom_point(aes(y=CumMExcDeathPer)) +
23   geom_errorbar(aes(ymin=CumLwExcDeathPer, ymax=CumUpExcDeathPer))+ theme_bw()+

```



```

24     labs(x='CovidCumDeath', y='Monthly CumDeath')
25     ggsave(file.path(out.dir, "MycityCum.png"),p,w =10,h = 5)
26     include_graphics(file.path(out.dir, "MycityCum.png"))

```

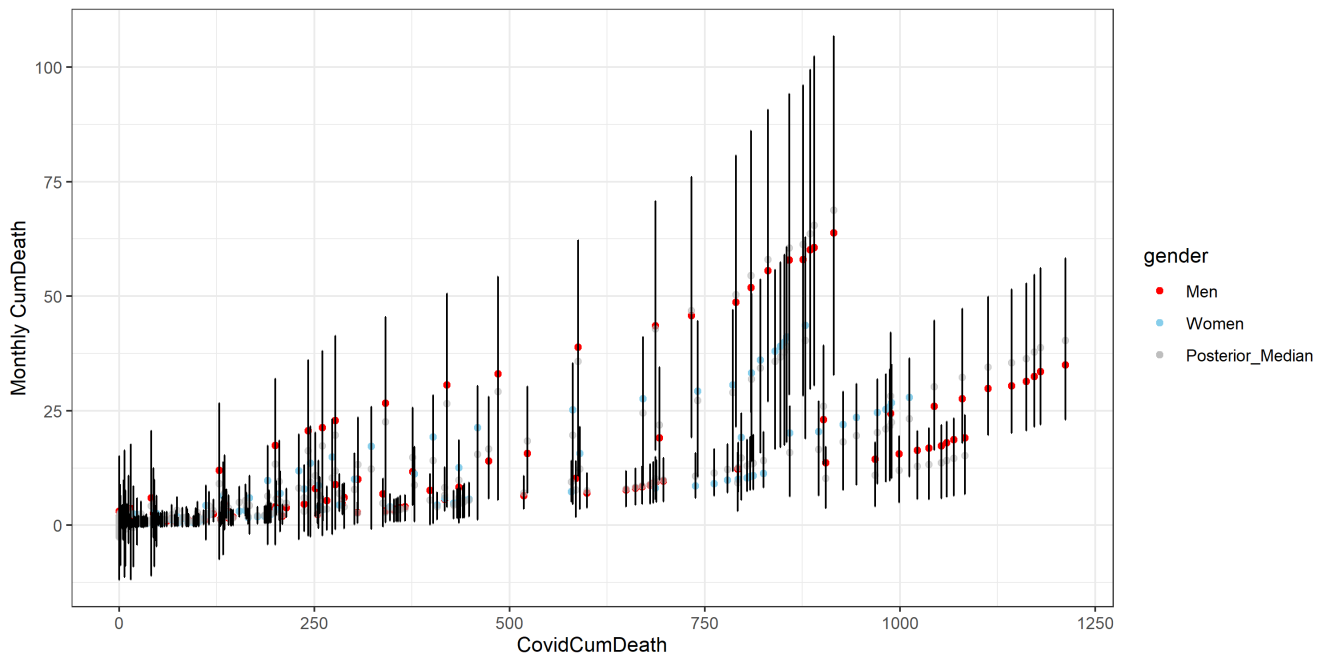


Figure 12: The cumulating monthly excess deaths per 1,000 population in each gender and age group against the corresponding cumulating monthly reported COVID-19 deaths with posterior medians and 95% credible intervals for the estimated cumulating monthly excess deaths per 1,000 population. Red points represent the real cumulating monthly excess deaths per 1,000 population for men and the blue points represent that for women and the half-transparent black points are the posterior medians

7.3 Calculate the difference in total excess deaths and the reported COVID-19 deaths by the end of 2021 in the city that you are assigned. Report posterior median estimates along with 95% posterior credible intervals for both men and women in a table.

```

1 Mycity3 <- Mycity2[, .(CovidDeaths = sum(CovidDeaths) ,
2                       exc_deaths = sum(exc_deaths)), by = c('year'))
3 Mycity3[,DIF := exc_deaths-CovidDeaths]
4 sum(Mycity3[which(year <= 2021),'DIF'])
5 # 7469

```

	year	CovidDeaths	exc_deaths	DIF
1	2019	0	0.00	0.00
2	2020	2324	5706.00	3382.00
3	2021	5485	9572.00	4087.00
4	2022	160	605.00	445.00

Table 4: The Data in *Mycity3*

Thus, by the end of 2021 in the city Belo Horizontewe, the total difference is 4087.

```

1 Mycity4 <- Mycity2[,.(PP_CL = sum(PP_CL),
2                       PP_CU = sum(PP_CU),
3                       PP_M = sum(PP_M)), by = c('gender')]

```

The posterior median estimates along with 95% posterior credible intervals for both men and women is shown in the following table 5

	gender	Lower Bound	Upper Bound	Median
1	Men	22489	32294	27223
2	Women	19222	28170	23522

Table 5: The posterior median estimates along with 95% posterior credible intervals for both men and women

8 Evaluating prediction accuracy

8.1

Add prediction and log like to the best model and get the model fit. (Inherit the stan data from Q6)

```

1 poi_regression3_txt <- "
2 data{
3   int<lower=1> N;
4   int<lower=1> K;
5   int<lower=0> y[N];
6   vector[N] offset;
7   matrix[N,K] X;
8 }
9 parameters{
10   real beta0;
11   vector[K] beta;
12 }
13 transformed parameters{
14   vector[N] log_lambda;
15   log_lambda = beta0 + X * beta + offset;
16 }
17 model{
18   beta0 ~ normal( 0 , 10 );
19   beta ~ normal( 0 , 1 );
20   y ~ poisson_log( log_lambda );
21 }
22
23 generated quantities {
24   int ypred[N];
25   real log_lik[N];
26
27   ypred = poisson_log_rng(beta0 + X * beta + offset);
28   for(i in 1:N)
29   {
30     log_lik[i] = poisson_log_lpmf(y[i] | beta0 + X[i,] * beta + offset);
31   }
32 }
33 "
34 if (!file.exists(file.path(out.dir, paste0('poi_deaths_fit3_',city_of_student_ascii,'.rds'))))
35 {
36   poi_regression3_compiled <- rstan::stan_model(
37     model_name = 'poi_regression3',

```

```

38   model_code = gsub('\t', ' ', poi_regression3_txt)
39   )
40   #m1_init_beta0 <- subset(dds, age_group_id == 1 & month_id == 1)[, log(deaths/population)]
41   m3_fit <- rstan::sampling(poi_regression3_compiled,
42     data = stan_data,
43     warmup = 5e2, iter = 5e3, chains = 4,
44     init = list(list(beta0 = 0 , beta = rep(0, stan_data$K)),
45                 list(beta0 = 0 , beta = rep(0, stan_data$K)),
46                 list(beta0 = 0 , beta = rep(0, stan_data$K)),
47                 list(beta0 = 0 , beta = rep(0, stan_data$K))
48               )
49   )
50   saveRDS(m3_fit, file = file.path(out.dir, paste0('poi_deaths_fit3_', city_of_student_ascii, '.rds')))
51 }
52 if (file.exists(file.path(out.dir, paste0('poi_deaths_fit3_', city_of_student_ascii, '.rds'))))
53 {
54   m3_fit <- readRDS( file.path(out.dir, paste0('poi_deaths_fit3_', city_of_student_ascii, '.rds')) )
55 }

```

MEAN ABSOLUTE ERROR and LOG POINTWISE PREDICTIVE DENSITY code::

```

1  est_para <- rstan::summary(m3_fit)$summary
2
3  log_like_est <- est_para[grepl('log_lik', rownames(est_para))][,1]
4  ypred_est <- est_para[grepl('ypred', rownames(est_para))][,1]
5  Mycity5 <- cbind(Mycity2, log_like_est, ypred_est)
6
7  Mycity5[, Ae:= abs(deaths-ypred_est)]
8  #Report
9  Mae <- Mycity5[, c('Ae'), with = FALSE] %>% as.vector()
10 Mae_val <- sum(Mae)/518
11 #63.60698
12
13 # Expected log pointwise predictive density of the 518 all-cause death values
14 exppd <- sum(log_like_est)
15 #-208236663
16
17 # Expected log pointwise predictive density of the all-cause death values that correspond to the last month in 2021.
18 exppd_36 <- Mycity5[which(month_id == 36), c('log_like_est'), with = FALSE]
19 exppd_36 <- sum(exppd_36)
20 # -1075481

```

Denote by z_i a new data point, and by z_i^* the corresponding, predicted pointwise estimate under a fitted model. Then the MEAN ABSOLUTE ERROR is

$$MAE = \frac{1}{n^*} \sum_{i=1}^{n^*} (|z_i - z_i^*|)$$

Then we calculate and report the mean absolute error between the actual 518 all-cause death values and the corresponding posterior median estimates: (Here we take z_i to be all-cause death and z_i^* to be the corresponding posterior median estimates) The mean absolute error between the actual 518 all-cause death values and the corresponding posterior median estimate is 63.60698.

Then we apply the LOG POINTWISE PREDICTIVE DENSITY which is numerically estimated with

$$\text{lppd.est} = \frac{1}{S} \sum_{s=1}^S \sum_{i=1}^{n^*} \log p(z_i | \theta^s)$$

The expected log point-wise predictive density of the 518 all-cause death values is -208236663 and the expected log point-wise predictive density of the all-cause death values that correspond to the last month in 2021 is -1075481.

8.2

```
1 # calculate approximate LOO cross validation for model without variable selection
2 m3_fit_loo <- loo::loo(m3_fit)
3 m3_fit_loo
4 # Computed from 18000 by 518 log-likelihood matrix
5 #
6 # -----
7 # Monte Carlo SE of elpd_loo is NA.
8 #
9 # Pareto k diagnostic values:
10 #               Count Pct.    Min. n_eff
11 # (-Inf, 0.5] (good)     0    0.0% <NA>
12 # (0.5, 0.7] (ok)       0    0.0% <NA>
13 # (0.7, 1] (bad)        0    0.0% <NA>
14 # (1, Inf) (very bad) 518 100.0% 0
15 # See help('pareto-k-diagnostic') for details.
16 m3_fit_loo$elpd_loo
17 # [1] -43025176
```

We see that the leave-one-out expected log point-wise predictive density is -43025176.

This numeric approximation is non-stable and high Pareto k diagnostic values. Vehtari (2016) suggest to refit the model to y_{-1} when Pareto k parameter $\hat{\zeta} > 0.7$ for some y_i . Re-fit your model Command::

```
1 dpop <- data
2 #data populaiton & deaths
3 dds <- merge(dd,dpop, by = c("city",'age_group', 'gender'))
4 city_of_student_ascii ='Belo Horizonte'
5 dds <- dds[which(city == 'Belo Horizonte',month_id=='36' ),]
6
7 design_matrix_X <- model.matrix(~ -1 + age_group+ gender, data = dds)
8 design_matrix_X <- design_matrix_X[,-1]
9
10 stan_data <- list()
11 stan_data$N <- nrow(dds)
12 stan_data$K <- ncol(design_matrix_X )
13 stan_data$y <- dds$deaths
14 stan_data$offset <- log(dds$population)
15 stan_data$X <- design_matrix_X
16
17 m4_fit <- rstan::sampling(poi_regression3_compiled,
18   data = stan_data,
19   warmup = 5e2, iter = 5e3, chains = 4,
20   init = list(list(beta0 = 0 , beta = rep(0, stan_data$K)),
21     list(beta0 = 0 , beta = rep(0, stan_data$K)),
22     list(beta0 = 0 , beta = rep(0, stan_data$K)),
23     list(beta0 = 0 , beta = rep(0, stan_data$K))
24   )
25 )
26 saveRDS(m4_fit, file = file.path(out.dir, paste0('poi_deaths_fit4_',city_of_student_ascii,'.rds')))
27 if (file.exists(file.path(out.dir, paste0('poi_deaths_fit4_',city_of_student_ascii,'.rds'))))
28 {
29   m4_fit <- readRDS( file.path(out.dir, paste0('poi_deaths_fit4_',city_of_student_ascii,'.rds')) )
30 }
31
32 est_para <- rstan::summary(m4_fit_loo)
33 log_like_est <- est_para[grepl('log_lik',rownames(est_para)),][,1]
34 ypred_est <- est_para[grepl('ypred',rownames(est_para)),][,1]
```

```

35
36 MAE <- sum(abs(dds$deaths - ypred_est))/518
37 #34.26034
38
39 exp_poster_density <- sum(log_like_est)
40 #[1] -37045684

```

For the last month in 2021, we have the Mean absolute value 34.26034 and expected log pointwise predictive density of the data -37045684.

Denote by z_i all-cause death, and by z_i^* the corresponding posterior median estimates under a fitted model. Then the MEAN ABSOLUTE ERROR is

$$MAE = \frac{1}{n^*} \sum_{i=1}^{n^*} (|z_i - z_i^*|)$$

Then we calculate and report the mean absolute error between the actual 518 all-cause death values and the corresponding posterior median estimates.

The LOG POINTWISE PREDICTIVE DENSITY (within-sample predictive accuracy) is

$$\text{lppd} = \log \prod_{i=1}^{n^*} p(y_i | y) = \log \prod_{i=1}^{n^*} \int p(y_i | \theta) p(\theta | y) d\theta$$

and the expected LOG POINTWISE PREDICTIVE DENSITY is

$$\text{lppd.est} = \frac{1}{S} \sum_{s=1}^S \sum_{i=1}^{n^*} \log p(z_i | \theta^s)$$

8.3 Rio de Janeiro

Inherit the stan model *poi_regression3.txt* from question 8.1.

```

1 dpop <- data
2 #data populaiton & deaths
3 dds <- merge(dd,dpop, by = c("city", 'age_group', 'gender'))
4 city_of_student_ascii = 'Belo Horizonte'
5 dds <- dds[which(city == 'Rio de Janeiro'),]
6
7 design_matrix_X <- model.matrix(~ -1 + age_group+ gender + as.factor(month_id), data = dds)
8 design_matrix_X <- design_matrix_X[,-1]
9
10 stan_data <- list()
11 stan_data$N <- nrow(dds)
12 stan_data$K <- ncol(design_matrix_X )
13 stan_data$y <- dds$deaths
14 stan_data$offset <- log(dds$population)
15 stan_data$X <- design_matrix_X
16 RJ_fit <- rstan::sampling(poi_regression3_compiled,
17   data = stan_data,
18   warmup = 5e2, iter = 5e3, chains = 4,
19   init = list(list(beta0 = 0 , beta = rep(0, stan_data$K)),
20     list(beta0 = 0 , beta = rep(0, stan_data$K)),
21     list(beta0 = 0 , beta = rep(0, stan_data$K)),
22     list(beta0 = 0 , beta = rep(0, stan_data$K))
23   )
24 )
25
26 saveRDS(RJ_fit, file.path(out.dir, "RJ_fit.rds"))
27

```

```

28 est_para_RJ <- rstan::summary(RJ_fit)$summary
29 ypred_RJ <- est_para_RJ[which(grepl('ypred', rownames(est_para_RJ))),]
30 loglike_RJ <- est_para_RJ[which(grepl('log_lik', rownames(est_para_RJ))),]
31 RJ_fit_loo <- loo::loo(RJ_fit)
32
33
34 yval <- cbind(ypred_RJ, dds) %>% as.data.table()
35 yval[,id:= 1:nrow(yval)]
36 # need cumulating
37 setkey(yval , 'month_id')
38 yCummu <- yval[, .(Cumdeaths = cumsum(deaths),
39                      CumMean = cumsum(mean),
40                      CumCL = cumsum(`2.5%`),
41                      CumCU = cumsum(`97.5%`),
42                      month_id = month_id), by = .(age_group, gender))]
43
44 p <-ggplot(yCummu, aes(x =month_id ))+ geom_point(aes(y = Cumdeaths), col = 'red') +
45   geom_point(aes(y = CumMean),alpha = 0.3)+
46   geom_errorbar(aes(ymin = CumCL, ymax = CumCU),alpha = 0.5)+
47   theme_bw() +
48   labs(x='month_id', y='posterior predictive vs actual deaths\n') +
49   facet_grid(gender~age_group)+
50   coord_cartesian(ylim = c(0,30000),xlim = c(0,37))
51
52 ggsave(file.path(out.dir, 'RJdeathCum.png'),p,w=10,h =5)
53 include_graphics(file.path(out.dir, 'RJdeathCum.png'))
54
55
56 Mae_RJ<- sum(abs(yval$mean - yval$deaths))/nrow(dds)
57 exp_pos_density <- sum(loglike_RJ[,1])
58 cat('MAE Value: ', Mae_RJ, "\n Expected Posterior Density: ", exp_pos_density)
59 # MAE Value: 29.92225
60 # Expected Posterior Density: -126470044

```

Fig 13 shows the comparison for the cumulating monthly all-cause death and the predicted cumulating monthly all-cause death and Fig 14 shows the comparison for the monthly all-cause death and the predicted monthly all-cause death.

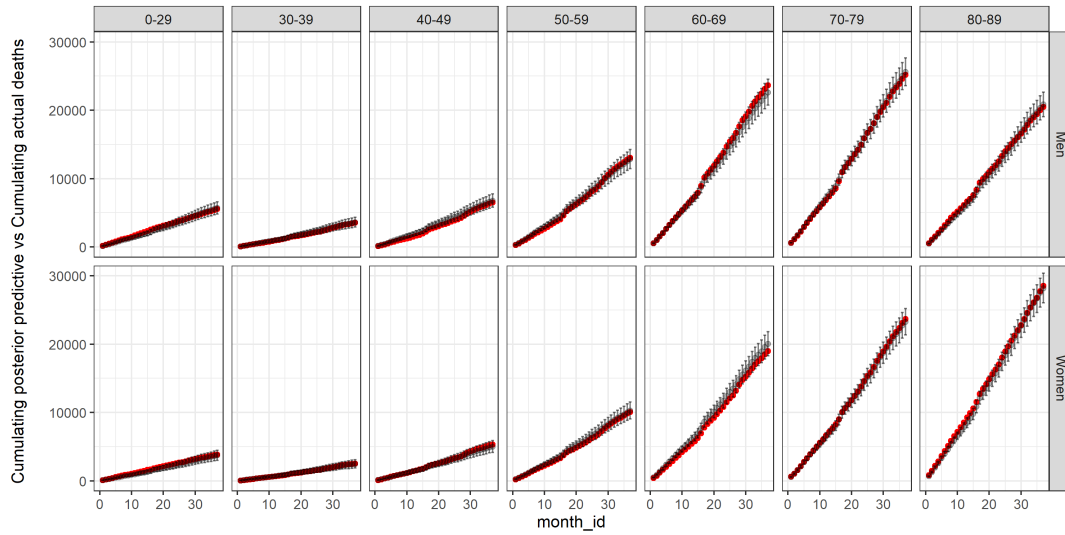


Figure 13: The plot for comparing the predictions to the actual data. The red points represent the actual cumulating monthly all-cause death and the black points represent the predicted cumulating monthly all-cause death. The 95% posterior credible intervals are also included in the graph.

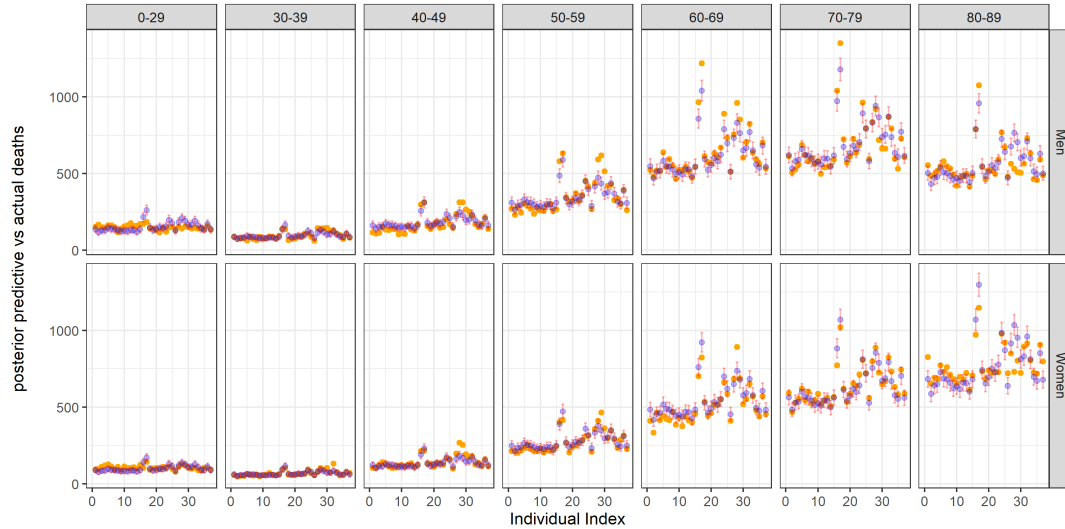


Figure 14: The plot for comparing the predictions to the actual data. The red points represent the actual monthly all-cause death and the blue points represent the predicted monthly all-cause death. The 95% posterior credible intervals are also included in the graph.

The mean absolute error between the actual 518 all-cause death values and the corresponding posterior mean estimates is 29.92225, and the expected log pointwise predictive density is -126470044.

8.4 Comment on the within-sample and out-of-sample prediction accuracies of your best model.

```

1 #For My city Belo Horizonte
2 m4_fit_loo <- loo::loo(m4_fit)
3 m4_fit_loo$estimates
4 #expected log pointwise predictive density
5 m4_fit_loo$estimates[1,1]
6 # [1] -40406698
7

```

```

8 exp_poster_density -m4_fit_loo$estimates[1,1]
9 #[1] 3361014
10
11 #For RJ
12 pred_density_loo <-RJ_fit_loo$elpd_loo
13 #[1] -136935950
14 bay_free_para <- exp_pos_density - pred_density_loo
15 #[1] 10465906

```

We see that the within-sample accuracy would give smaller prediction accuracy than the out-of-sample accuracy. By having the within-sample and out-of-sample prediction accuracies, we can calculate the BAYESIAN LOO FREE PARAMETERS of a model which is the difference between within-sample accuracy and the out-of-sample accuracy,

$$p.loo.est = lppd.est - lppd.loo.cv.est$$

measuring of model complexity, and then make model comparison with other models. For city Belo Horizonte, the BAYESIAN LOO FREE PARAMETERS are 3361014 and for Rio de Janeiro is 10465906. We see the within-sample accuracy for this model highly overestimates the out-of-sample predictive accuracy. The higher *lppd.loo.cv.est*, the higher predictive accuracy the model is.